

Feature Detection and Matching: Part 1

Chul Min Yeum

Assistant Professor

Civil and Environmental Engineering

University of Waterloo, Canada

CIVE 497 – CIVE 700: Smart Structure Technology

Last updated: 2021-02-22



UNIVERSITY OF WATERLOO
FACULTY OF ENGINEERING

Two Major Challenges in Computer Vision

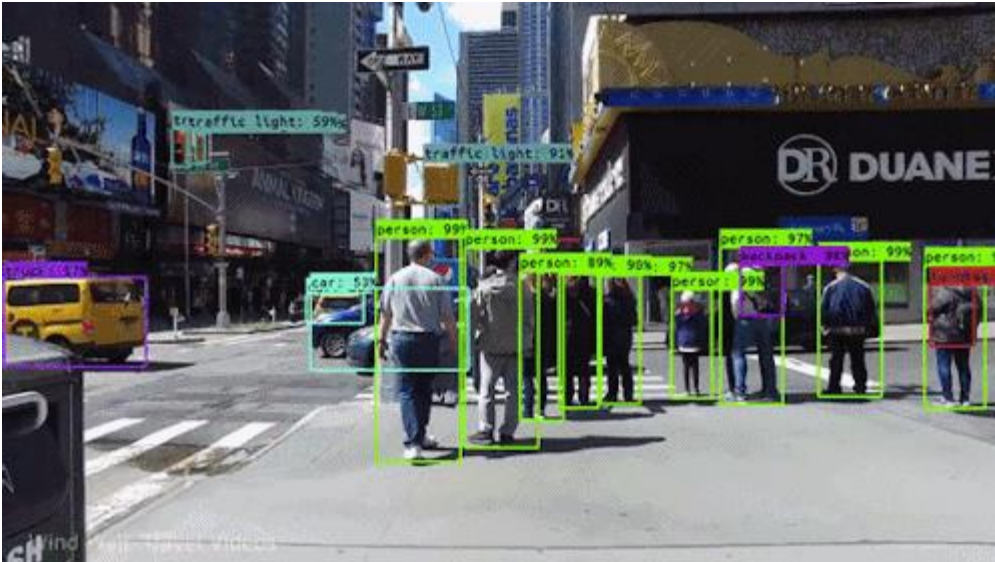
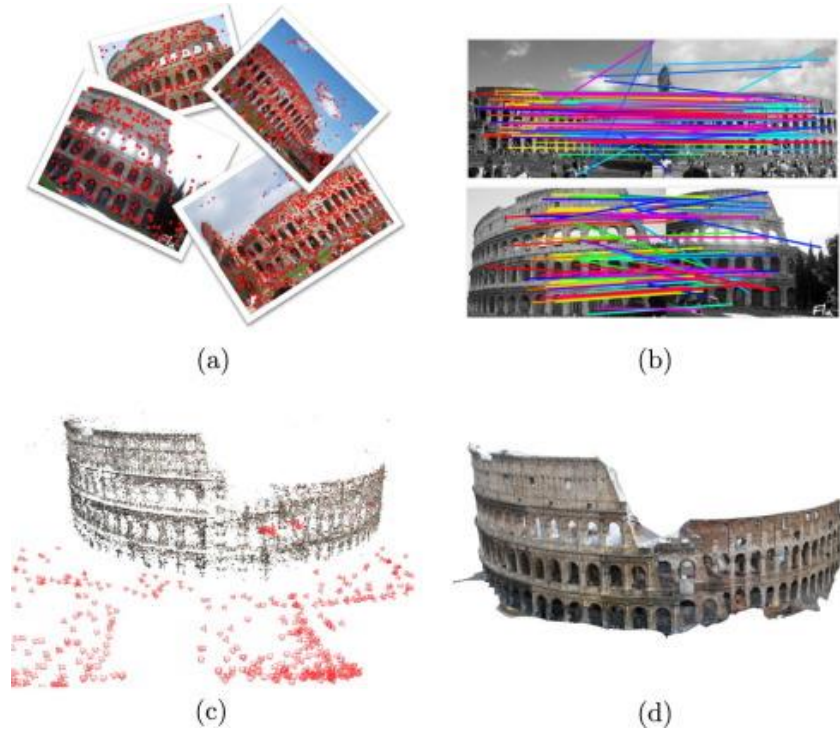


Image recognition (interpretation)




Feature matching (3D reconstruction)

Why Do We Extract Features?

- We have two images
- How do we combine or stitch them together?



Why We Extract Features? Extracting Features (Continue)



What are Features (Keypoints)?

- A feature is a prominent point that is selected based on a certain criteria, such as edge, corner, or blob.
- This is represented in terms of the coordinates of the image points by pixel or sub-pixel.
- The feature likely contain and preserve the distinctive **local** regional information.
- Note: “interest points” = “keypoints”, also sometimes called “features”

Many applications:

- Object/motion tracking: which points are good to track?
- 3D scene reconstruction: find correspondences across different views
- Object recognition: find patches likely to tell us something about object category

Example: Automatic Panoramas



Structure from motion software (Pix 4D)



A large volume collected images from drone



Orthophoto($10,000 \times 3,656$) geometrically connected to each collected images


Jongseong Choi, Chul Min Yeum, Shirley J. Dyke, and Mohammad R. Jahanshahi, "Computer-Aided Approach for Rapid Post-Event Visual Evaluation of a Building Façade," *Sensors*, 18, 3017 (2018).

Image or Feature Matching Is a Challenging Task




Example: Keypoints/Features

Q. Suppose you have to click on some point to find out how the image is deformed. Which points would you choose?

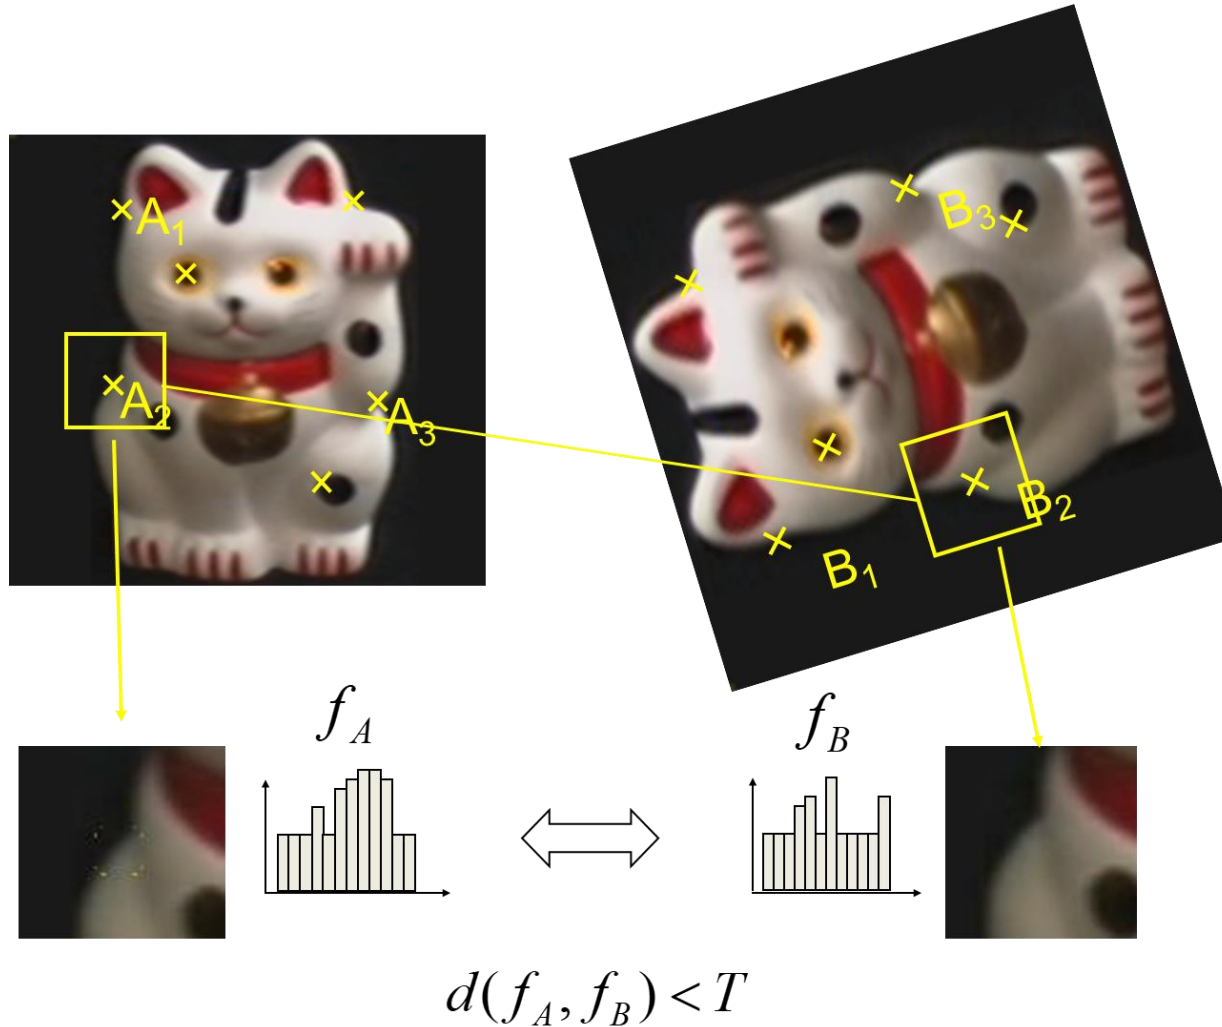


original



deformed

Overview of Feature Matching



Step 1. Find a set of distinctive key-points

Step 2. Define a region around each keypoint

Step 3. Extract and normalize the region content

Step 4. Compute a local descriptor from the normalized region

Step 5. Match local descriptors

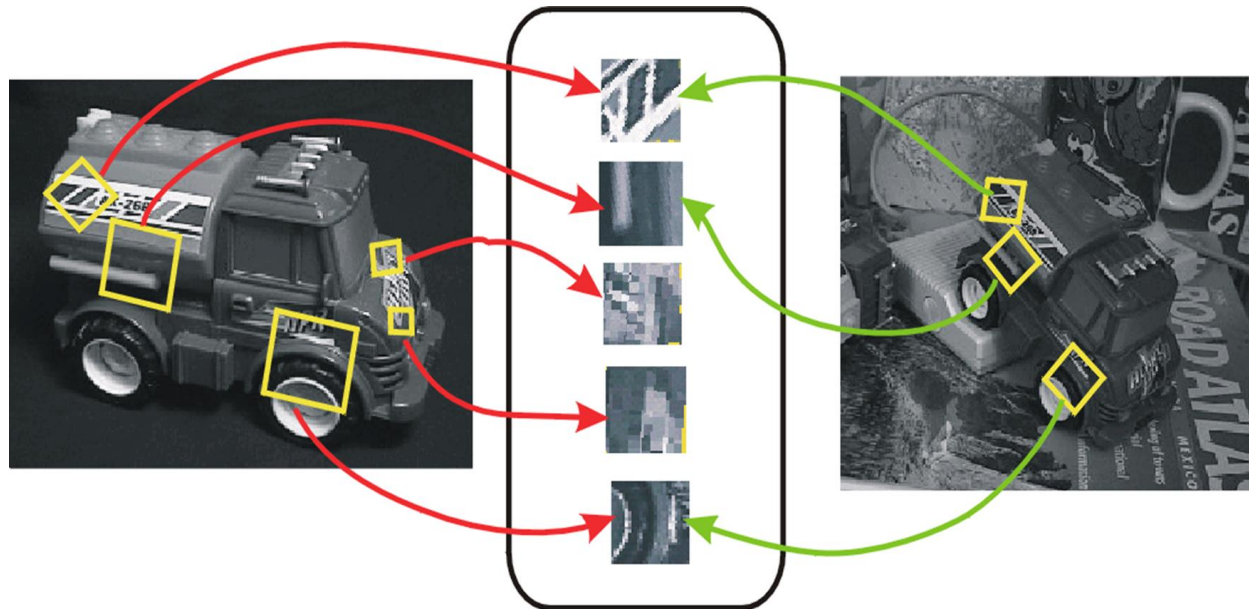
Characteristic of Good Features

Repeatability: The same feature can be found in several images despite geometric and photometric transformations



Saliency: Each feature is distinctive

Compactness and efficiency: Many fewer features than image pixels

Locality: A feature occupies a relatively small area of the image; robust to clutter and occlusion



Example: Step 1. Find a Set of Distinctive Keypoints



Good keypoints

Not good keypoints

Step 1. Find a set of distinctive keypoints

Step 2. Define a region around each keypoint

Step 3. Extract and normalize the region content

Step 4. Compute a local descriptor from the normalized region

Step 5. Match local descriptors

Q. How to match the extracted features?

Example: Step 2. Define a Region Around Each Keypoint



How to describe the features?



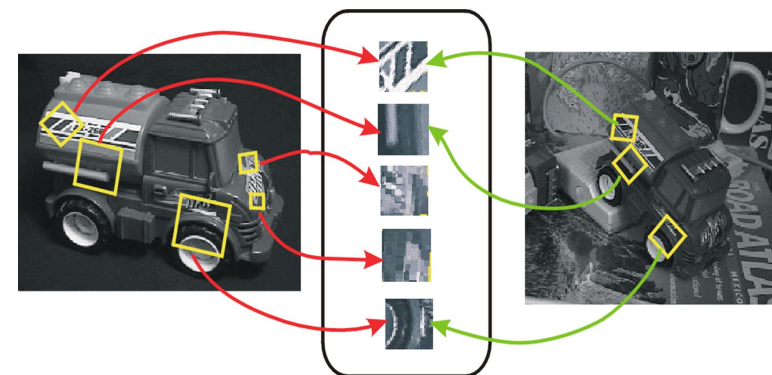
Step 1. Find a set of distinctive keypoints

Step 2. Define a region around each keypoint

Step 3. Extract and normalize the region content

Step 4. Compute a local descriptor from the normalized region


Step 5. Match local descriptors



Example: Step 3. Extract and Normalize the Region Content



How to make the “region”
insensitive to color or light variations



Step 1. Find a set of distinctive key-points

Step 2. Define a region around each keypoint

Step 3. Extract and normalize the region content


Step 4. Compute a local descriptor from the normalized region

Step 5. Match local descriptors

Example: Step 4. Compute a Local Descriptor from the Normalized Region



How to represent the region using a fixed size vector



10
8
⋮
5
6

4
2
⋮
9
10

10
8
⋮
4
6

4
2
⋮
9
11

Step 1. Find a set of distinctive key-points

Step 2. Define a region around each keypoint


Step 3. Extract and normalize the region content

Step 4. Compute a local descriptor from the normalized region


Step 5. Match local descriptors

The descriptor is a fingerprint that can uniquely represent the features.

Example: Step 5. Match Local Descriptors



Matching the descriptors



Step 1. Find a set of distinctive key-points

Step 2. Define a region around each keypoint

Step 3. Extract and normalize the region content

Step 4. Compute a local descriptor from the normalized region


Step 5. Match local descriptors

What Locations Would You Choose as Features ?




Corner Detection: Basic Idea

- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



Flat region


no change in all directions



Edge

no change along the edge
direction

Distinctive but not locality



Corner


significant change in all
directions

Distinctive and locality


Invariance and Covariance of Features

Invariance: images are transformed, and feature locations do not change


Covariance: if we have two transformed versions of the same image, features should be detected in corresponding locations




Original



Translation



Rotation




Intensity







Scale

Intensity Change




$$I \rightarrow aI + b$$

- Only derivatives are used for edge or corner detection \Rightarrow invariance to intensity shift $I \rightarrow I + b$
- Intensity scaling: $I \rightarrow aI$



Partially invariant to intensity change


Image Translation



- Derivatives and window function are shift-invariant

Corner locations are covariant to image translation


Image Rotation






Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner locations are covariant to image rotation

Scaling



Corner




All points will be classified as **edges**

Corner location is **NOT** invariant to image scale!


Scale Invariant Detection

- Consider regions (e.g. circles) of different sizes around a point
- Regions of corresponding sizes will be the same in both images



Scale Invariant Detection (Continue)

- To detect the corners in both image, we need to know how to choose corresponding circles **independently** in each image?
- Choose the scale of the “best” corner



Example: Scale Invariance



Slide Credits and References

- Lecture notes: Gordon Wetzstein
- Lecture notes: Mohammad Jahanshahi
- Lecture notes: Noah Snavely
- Lecture notes: L. Fei-Fei
- Lecture notes: D. Forsyth
- Lecture notes: James Hayes
- Lecture notes: Yacov Hel-Or
- Lecture notes: K. Grauman, B. Leibe